



МИНОБРНАУКИ РОССИИ
Федеральное государственное бюджетное образовательное учреждение высшего
образования
«МИРЭА – Российский технологический университет»
РТУ МИРЭА

ИКБ направление «Киберразведка и противодействие угрозам с
применением технологий искусственного интеллекта» 10.04.01

Кафедра КБ-4 «Интеллектуальные системы информационной
безопасности»

Лабораторная работа №1

по дисциплине

«Анализ защищенности систем искусственного интеллекта»

Группа:
ББМО-01-22
Выполнил:
Орлянский В.Е.

Проверил:
Спирин А.А.

Москва 2023

1. Клонировать репозиторий.

```
!git clone https://github.com/ewatson2/EEL6812_DeepFool_Project.git
```

2. Сменим директорию исполнения на вновь созданную папку "EEL6812_DeepFool_Project" проекта.

```
%cd /content/EEL6812_DeepFool_Project
```

3. Выполним импорт стандартных и вспомогательных библиотек.

```
import numpy as np
import os
import json, torch
from torch.utils.data import DataLoader, random_split
from torchvision import datasets, models
from torchvision.transforms import transforms
from models.project_models import FC_500_150, LeNet_CIFAR, LeNet_MNIST, Net
from utils.project_utils import get_clip_bounds, evaluate_attack, display_attack
```

4. Установим случайное рандомное значение на номер в списке (38).

```
rand_seed = 38
np.random.seed(rand_seed)
torch.manual_seed(rand_seed)

use_cuda = torch.cuda.is_available()
device = torch.device('cuda' if use_cuda else 'cpu')
```

5. Загрузим датасет MNIST.

```
mnist_mean = 0.5
mnist_std = 0.5
mnist_dim = 28

mnist_min, mnist_max = get_clip_bounds(mnist_mean, mnist_std, mnist_dim)
mnist_min = mnist_min.to(device)
mnist_max = mnist_max.to(device)

mnist_tf = transforms.Compose([ transforms.ToTensor(), transforms.Normalize( mean=mnist_mean, std=mnist_std)])

mnist_tf_train = transforms.Compose([ transforms.RandomHorizontalFlip(), transforms.ToTensor(), transforms.Normalize( mean=mnist_mean, std=mnist_std)])

mnist_tf_inv = transforms.Compose([ transforms.Normalize( mean=0.0, std=np.divide(1.0, mnist_std)), transforms.Normalize( mean=np.multiply(-1.0, mnist_std), std=1.0)])

mnist_temp = datasets.MNIST(root='datasets/mnist', train=True, download=True, transform=mnist_tf_train)
mnist_train, mnist_val = random_split(mnist_temp, [50000, 10000])
mnist_test = datasets.MNIST(root='datasets/mnist', train=False, download=True, transform=mnist_tf)
```

6. Загрузим датасет CIFAR-10.

```
cifar_mean = [0.491, 0.482, 0.447]
cifar_std = [0.202, 0.199, 0.201]
cifar_dim = 32

cifar_min, cifar_max = get_clip_bounds(cifar_mean, cifar_std, cifar_dim)
cifar_min = cifar_min.to(device)
cifar_max = cifar_max.to(device)
cifar_tf = transforms.Compose([ transforms.ToTensor(), transforms.Normalize( mean=cifar_mean, std=cifar_std)])

cifar_tf_train = transforms.Compose([ transforms.RandomCrop( size=cifar_dim, padding=4), transforms.RandomHorizontalFlip(), transforms.ToTensor(), transforms.Normalize( mean=cifar_mean, std=cifar_std)])

cifar_tf_inv = transforms.Compose([ transforms.Normalize( mean=[0.0, 0.0, 0.0], std=np.divide(1.0, cifar_std)), transforms.Normalize( mean=np.multiply(-1.0, cifar_mean), std=[1.0, 1.0, 1.0])])

cifar_temp = datasets.CIFAR10(root='datasets/cifar-10', train=True, download=True, transform=cifar_tf_train)

cifar_train, cifar_val = random_split(cifar_temp, [40000, 10000])

cifar_test = datasets.CIFAR10(root='datasets/cifar-10', train=False, download=True, transform=cifar_tf)

cifar_classes = ['airplane', 'automobile', 'bird', 'cat', 'deer', 'dog', 'frog', 'horse', 'ship', 'truck']
```

7. Выполним настройку и загрузку DataLoader.

```
batch_size = 64
workers = 4
mnist_loader_train = DataLoader(mnist_train, batch_size=batch_size, shuffle=True, num_workers=workers)
mnist_loader_val = DataLoader(mnist_val, batch_size=batch_size, shuffle=False, num_workers=workers)
mnist_loader_test = DataLoader(mnist_test, batch_size=batch_size, shuffle=False, num_workers=workers)
cifar_loader_train = DataLoader(cifar_train, batch_size=batch_size, shuffle=True, num_workers=workers)
cifar_loader_val = DataLoader(cifar_val, batch_size=batch_size, shuffle=False, num_workers=workers)
cifar_loader_test = DataLoader(cifar_test, batch_size=batch_size, shuffle=False, num_workers=workers)
```

8. Настроим обучающую модель.

```
train_model = True

epochs = 50
epochs_nin = 100

lr = 0.004
lr_nin = 0.01
lr_scale = 0.5

momentum = 0.9

print_step = 5

deep_batch_size = 10
deep_num_classes = 10
deep_overshoot = 0.02
deep_max_iters = 50

deep_args = [deep_batch_size, deep_num_classes, deep_overshoot, deep_max_iters]

if not os.path.isdir('weights/deepfool'): os.makedirs('weights/deepfool', exist_ok=True)
if not os.path.isdir('weights/fgsm'): os.makedirs('weights/fgsm', exist_ok=True)
```

9. Загрузим и оценим стойкость модели Network-In-Network Model к FGSM и DeepFool атакам на основе датасета CIFAR-10.

```
fgsm_eps = 0.2
model = Net().to(device)
model.load_state_dict(torch.load('weights/clean/cifar_nin.pth', map_location=torch.device('cpu')))
evaluate_attack('cifar_nin_fgsm.csv', 'results', device, model, cifar_loader_test, cifar_min, cifar_max, fgsm_eps, is_fgsm=True)
print('')
evaluate_attack('cifar_nin_deepfool.csv', 'results', device, model, cifar_loader_test, cifar_min, cifar_max, deep_args, is_fgsm=False)
if device.type == 'cuda': torch.cuda.empty_cache()

FGSM Test Error : 81.29%
FGSM Robustness : 1.77e-01
FGSM Time (All Images) : 0.67 s
FGSM Time (Per Image) : 67.07 us

DeepFool Test Error : 93.76%
DeepFool Robustness : 2.12e-02
DeepFool Time (All Images) : 185.12 s
DeepFool Time (Per Image) : 18.51 ms
```

10. Загрузим и оценим стойкость модели LeNet к FGSM и DeepFool атакам на основе датасета CIFAR-10.

```

fgsm_eps = 0.1
model = LeNet_CIFAR().to(device)
model.load_state_dict(torch.load('weights/clean/cifar_lenet.pth', map_location=torch.device('cpu')))
evaluate_attack('cifar_lenet_fgsm.csv', 'results', device, model, cifar_loader_test, cifar_min, cifar_max, fgsm_eps, is_fgsm=True)
print('')
evaluate_attack('cifar_lenet_deepfool.csv', 'results', device, model, cifar_loader_test, cifar_min, cifar_max, deep_args, is_fgsm=False)
if device.type == 'cuda': torch.cuda.empty_cache()

FGSM Test Error : 91.71%
FGSM Robustness : 8.90e-02
FGSM Time (All Images) : 0.40 s
FGSM Time (Per Image) : 40.08 us

DeepFool Test Error : 87.81%
DeepFool Robustness : 1.78e-02
DeepFool Time (All Images) : 73.27 s
DeepFool Time (Per Image) : 7.33 ms

```

11. Выполним оценку атакующих примеров для сетей.

```

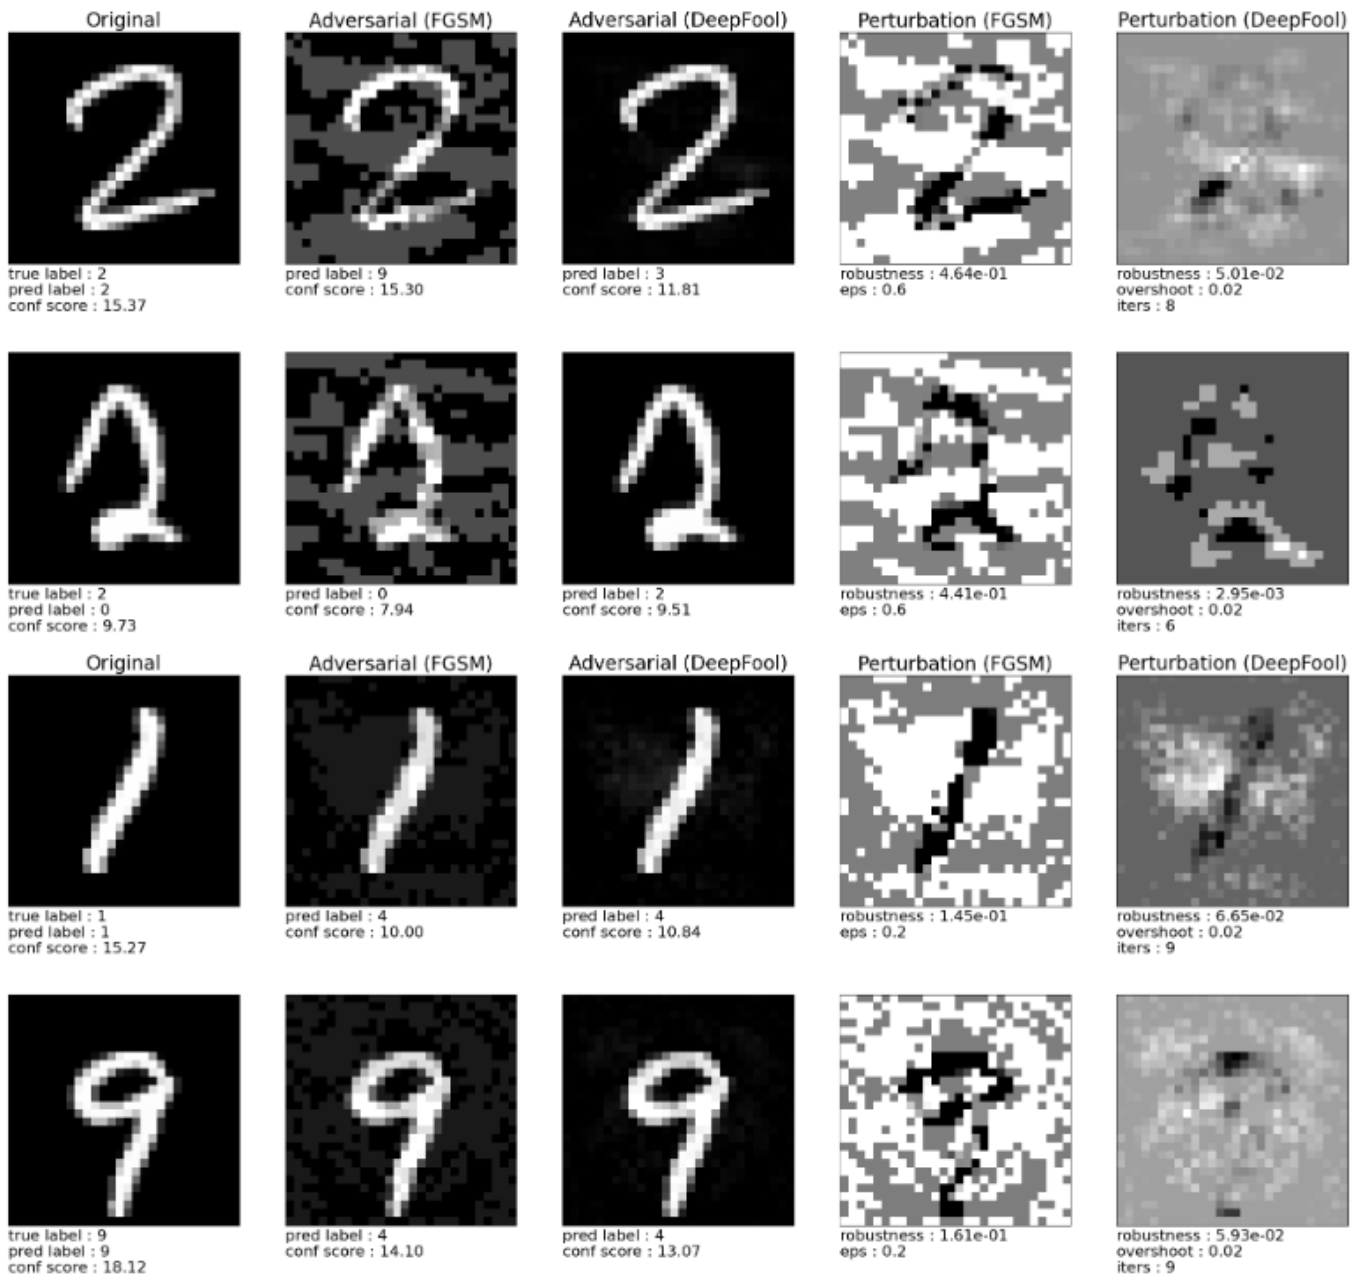
#LeNet
fgsm_eps = 0.6
model = LeNet_MNIST().to(device)
model.load_state_dict(torch.load('weights/clean/mnist_lenet.pth'))
display_attack(device, model, mnist_test, mnist_tf_inv, mnist_min, mnist_max, fgsm_eps, deep_args, has_
if device.type == 'cuda': torch.cuda.empty_cache()

#FCNet
fgsm_eps = 0.2
model = FC_500_150().to(device)
model.load_state_dict(torch.load('weights/clean/mnist_fc.pth'))
display_attack(device, model, mnist_test, mnist_tf_inv, mnist_min, mnist_max, fgsm_eps, deep_args, has_
if device.type == 'cuda': torch.cuda.empty_cache()

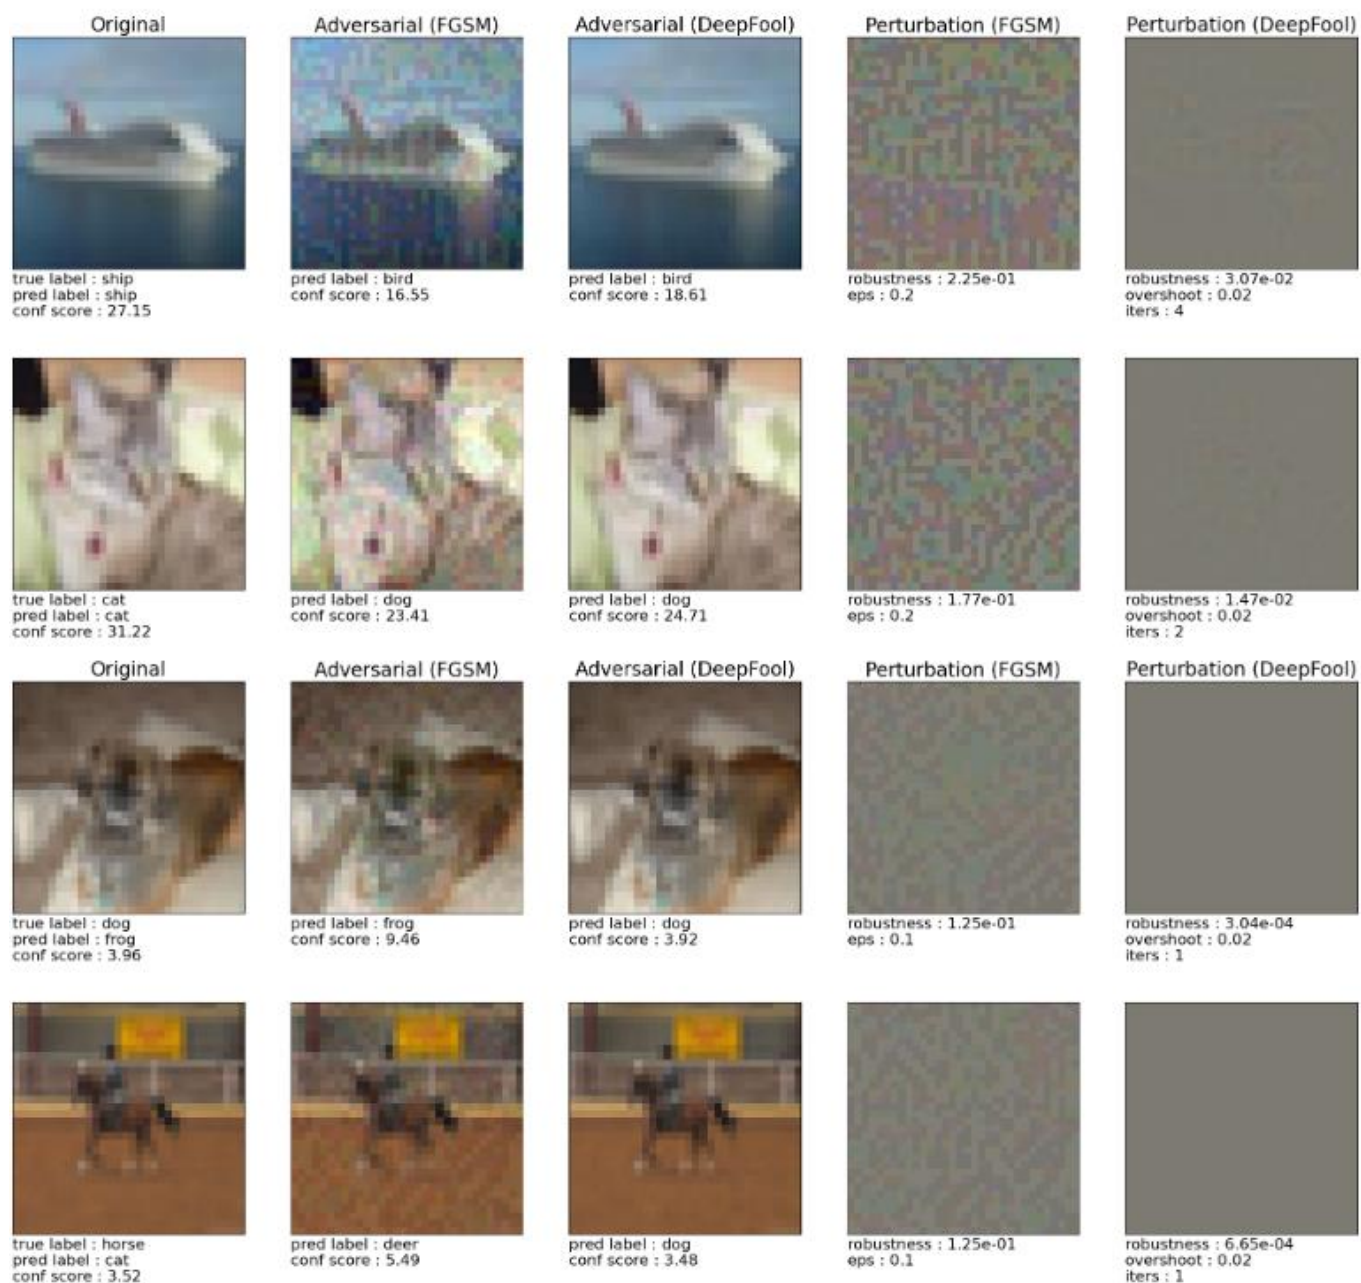
#Network-in-Network
fgsm_eps = 0.2
model = Net().to(device)
model.load_state_dict(torch.load('weights/clean/cifar_nin.pth'))
display_attack(device, model, cifar_test, cifar_tf_inv, cifar_min, cifar_max, fgsm_eps, deep_args, has_
if device.type == 'cuda': torch.cuda.empty_cache()

#LeNet CIFAR-10
fgsm_eps = 0.1
model = LeNet_CIFAR().to(device)
model.load_state_dict(torch.load('weights/clean/cifar_lenet.pth'))
display_attack(device, model, cifar_test, cifar_tf_inv, cifar_min, cifar_max, fgsm_eps, deep_args, has_
if device.type == 'cuda': torch.cuda.empty_cache()

```



\



12. Отражаем отличия для fgsm_eps=(0.001, 0.02, 0.5, 0.9, 10).

```

fgsm_eps = 0.001
model = FC_500_150().to(device)
model.load_state_dict(torch.load('weights/clean/mnist_fc.pth'))
display_attack(device, model, mnist_test, mnist_tf_inv, mnist_min, mnist_max, fgsm_eps, deep_args, has_
if device.type == 'cuda': torch.cuda.empty_cache()




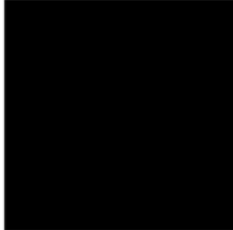
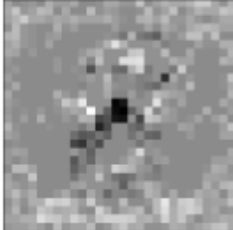




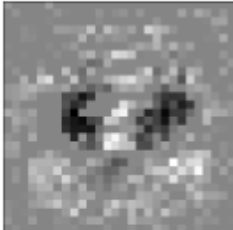




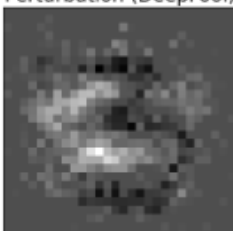



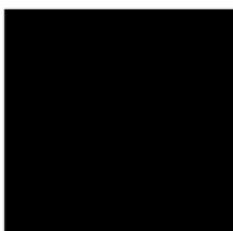

fgsm_eps = 0.002
model = FC_500_150().to(device)
model.load_state_dict(torch.load('weights/clean/mnist_fc.pth'))
display_attack(device, model, mnist_test, mnist_tf_inv, mnist_min, mnist_max, fgsm_eps, deep_args, has_
if device.type == 'cuda': torch.cuda.empty_cache()


fgsm_eps = 0.5
model = FC_500_150().to(device)
model.load_state_dict(torch.load('weights/clean/mnist_fc.pth'))
display_attack(device, model, mnist_test, mnist_tf_inv, mnist_min, mnist_max, fgsm_eps, deep_args, has_
if device.type == 'cuda': torch.cuda.empty_cache()

fgsm_eps = 0.9
model = FC_500_150().to(device)
model.load_state_dict(torch.load('weights/clean/mnist_fc.pth'))
display_attack(device, model, mnist_test, mnist_tf_inv, mnist_min, mnist_max, fgsm_eps, deep_args, has_
if device.type == 'cuda': torch.cuda.empty_cache()

fgsm_eps = 10
model = FC_500_150().to(device)
model.load_state_dict(torch.load('weights/clean/mnist_fc.pth'))
display_attack(device, model, mnist_test, mnist_tf_inv, mnist_min, mnist_max, fgsm_eps, deep_args, has_
if device.type == 'cuda': torch.cuda.empty_cache()

```

Original	Adversarial (FGSM)	Adversarial (DeepFool)	Perturbation (FGSM)	Perturbation (DeepFool)
 true label : 8 pred label : 8 conf score : 12.86	 pred label : 8 conf score : 12.79	 pred label : 9 conf score : 9.74	 robustness : 8.67e-04 eps : 0.001	 robustness : 4.25e-02 overshoot : 0.02 iters : 12
 true label : 9 pred label : 9 conf score : 20.62	 pred label : 9 conf score : 20.59	 pred label : 3 conf score : 13.74	 robustness : 8.58e-04 eps : 0.001	 robustness : 9.63e-02 overshoot : 0.02 iters : 10
 true label : 3 pred label : 3 conf score : 16.43	 pred label : 3 conf score : 16.32	 pred label : 8 conf score : 9.18	 robustness : 1.60e-03 eps : 0.002	 robustness : 6.94e-02 overshoot : 0.02 iters : 7
 true label : 0 pred label : 0 conf score : 20.63	 pred label : 0 conf score : 20.55	 pred label : 9 conf score : 12.24	 robustness : 1.71e-03 eps : 0.002	 robustness : 1.51e-01 overshoot : 0.02 iters : 10



13. Проверим влияние параметра fgsm_eps для FC на датасете MNIST.

```

fgsm_eps = 0.001
model = FC_500_150().to(device)
model.load_state_dict(torch.load('weights/clean/mnist_fc.pth'))
display_attack(device, model, mnist_test, mnist_tf_inv, mnist_min, mnist_max, fgsm_eps, deep_args, has_labels=False, :
if device.type == 'cuda': torch.cuda.empty_cache()





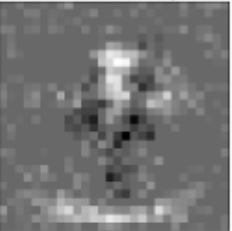



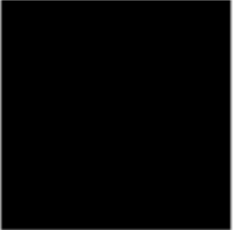
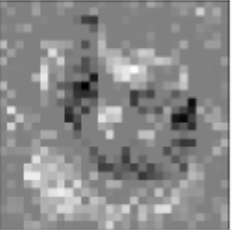




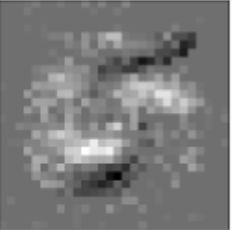




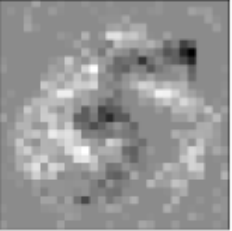
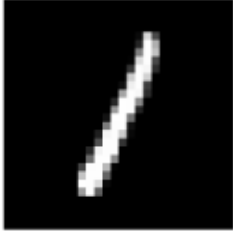



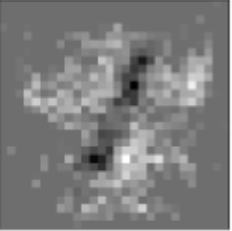
fgsm_eps = 0.02
model = FC_500_150().to(device)
model.load_state_dict(torch.load('weights/clean/mnist_fc.pth'))
display_attack(device, model, mnist_test, mnist_tf_inv, mnist_min, mnist_max, fgsm_eps, deep_args, has_labels=False, :
if device.type == 'cuda': torch.cuda.empty_cache()

fgsm_eps = 0.5
model = FC_500_150().to(device)
model.load_state_dict(torch.load('weights/clean/mnist_fc.pth'))
display_attack(device, model, mnist_test, mnist_tf_inv, mnist_min, mnist_max, fgsm_eps, deep_args, has_labels=False, :
if device.type == 'cuda': torch.cuda.empty_cache()

fgsm_eps = 0.9
model = FC_500_150().to(device)
model.load_state_dict(torch.load('weights/clean/mnist_fc.pth'))
display_attack(device, model, mnist_test, mnist_tf_inv, mnist_min, mnist_max, fgsm_eps, deep_args, has_labels=False, :
if device.type == 'cuda': torch.cuda.empty_cache()

fgsm_eps = 10
model = FC_500_150().to(device)
model.load_state_dict(torch.load('weights/clean/mnist_fc.pth'))
display_attack(device, model, mnist_test, mnist_tf_inv, mnist_min, mnist_max, fgsm_eps, deep_args, has_labels=False, :
if device.type == 'cuda': torch.cuda.empty_cache()

```

Original	Adversarial (FGSM)	Adversarial (DeepFool)	Perturbation (FGSM)	Perturbation (DeepFool)
 true label : 4 pred label : 4 conf score : 15.39	 pred label : 4 conf score : 15.35	 pred label : 9 conf score : 10.13	 robustness : 8.27e-04 eps : 0.001	 robustness : 8.53e-02 overshoot : 0.02 iters : 9
 true label : 6 pred label : 6 conf score : 19.03	 pred label : 6 conf score : 18.97	 pred label : 2 conf score : 9.49	 robustness : 7.79e-04 eps : 0.001	 robustness : 1.13e-01 overshoot : 0.02 iters : 12
 true label : 5 pred label : 5 conf score : 13.76	 pred label : 5 conf score : 12.40	 pred label : 8 conf score : 7.80	 robustness : 1.52e-02 eps : 0.02	 robustness : 5.22e-02 overshoot : 0.02 iters : 8
 true label : 5 pred label : 5 conf score : 20.37	 pred label : 5 conf score : 18.91	 pred label : 2 conf score : 10.55	 robustness : 1.61e-02 eps : 0.02	 robustness : 8.17e-02 overshoot : 0.02 iters : 11
 true label : 1 pred label : 1 conf score : 15.09	 pred label : 4 conf score : 6.44	 pred label : 8 conf score : 7.99	 robustness : 3.59e-01 eps : 0.5	 robustness : 7.63e-02 overshoot : 0.02 iters : 10



true label : 4
pred label : 4
conf score : 18.09

Original



pred label : 7
conf score : 20.71

Adversarial (FGSM)



pred label : 9
conf score : 10.21

Adversarial (DeepFool)




robustness : 4.08e-01
eps : 0.5

Perturbation (FGSM)



robustness : 1.05e-01
overshoot : 0.02
iters : 9

Perturbation (DeepFool)



true label : 9
pred label : 9
conf score : 13.50

Original



pred label : 4
conf score : 20.66

Adversarial (FGSM)




pred label : 4
conf score : 11.36

Adversarial (DeepFool)



robustness : 6.84e-01
eps : 0.9

Perturbation (FGSM)



robustness : 3.83e-02
overshoot : 0.02
iters : 10

Perturbation (DeepFool)



true label : 5
pred label : 5
conf score : 16.44

Original



pred label : 2
conf score : 14.07

Adversarial (FGSM)



pred label : 8
conf score : 8.73

Adversarial (DeepFool)




robustness : 6.79e-01
eps : 0.9

Perturbation (FGSM)



robustness : 6.39e-02
overshoot : 0.02
iters : 8

Perturbation (DeepFool)



true label : 3
pred label : 3
conf score : 22.63

Original




pred label : 5
conf score : 61.76

Adversarial (FGSM)




pred label : 5
conf score : 14.52

Adversarial (DeepFool)



robustness : 1.43e+00
eps : 10

Perturbation (FGSM)



robustness : 9.35e-02
overshoot : 0.02
iters : 10

Perturbation (DeepFool)



true label : 0
pred label : 0
conf score : 11.48

Original



pred label : 7
conf score : 56.61

Adversarial (FGSM)



pred label : 7
conf score : 8.74

Adversarial (DeepFool)



robustness : 1.56e+00
eps : 10

Perturbation (FGSM)



robustness : 4.94e-02
overshoot : 0.02
iters : 10

Perturbation (DeepFool)

14. Проверим влияние параметра fgsm_eps для LeNet на датасете MNIST.





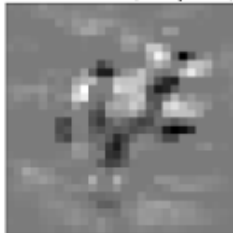




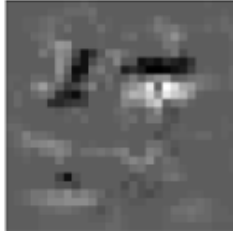
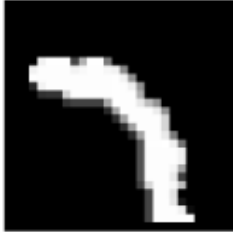



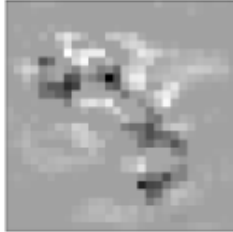










```
fgsm_eps = 0.001
model = LeNet_MNIST().to(device)
model.load_state_dict(torch.load('weights/clean/mnist_lenet.pth'))
display_attack(device, model, mnist_test, mnist_tf_inv, mnist_min, mnist_max, fgsm_eps, deep_args, has_labels=False, l2_norm=0)
if device.type == 'cuda': torch.cuda.empty_cache()

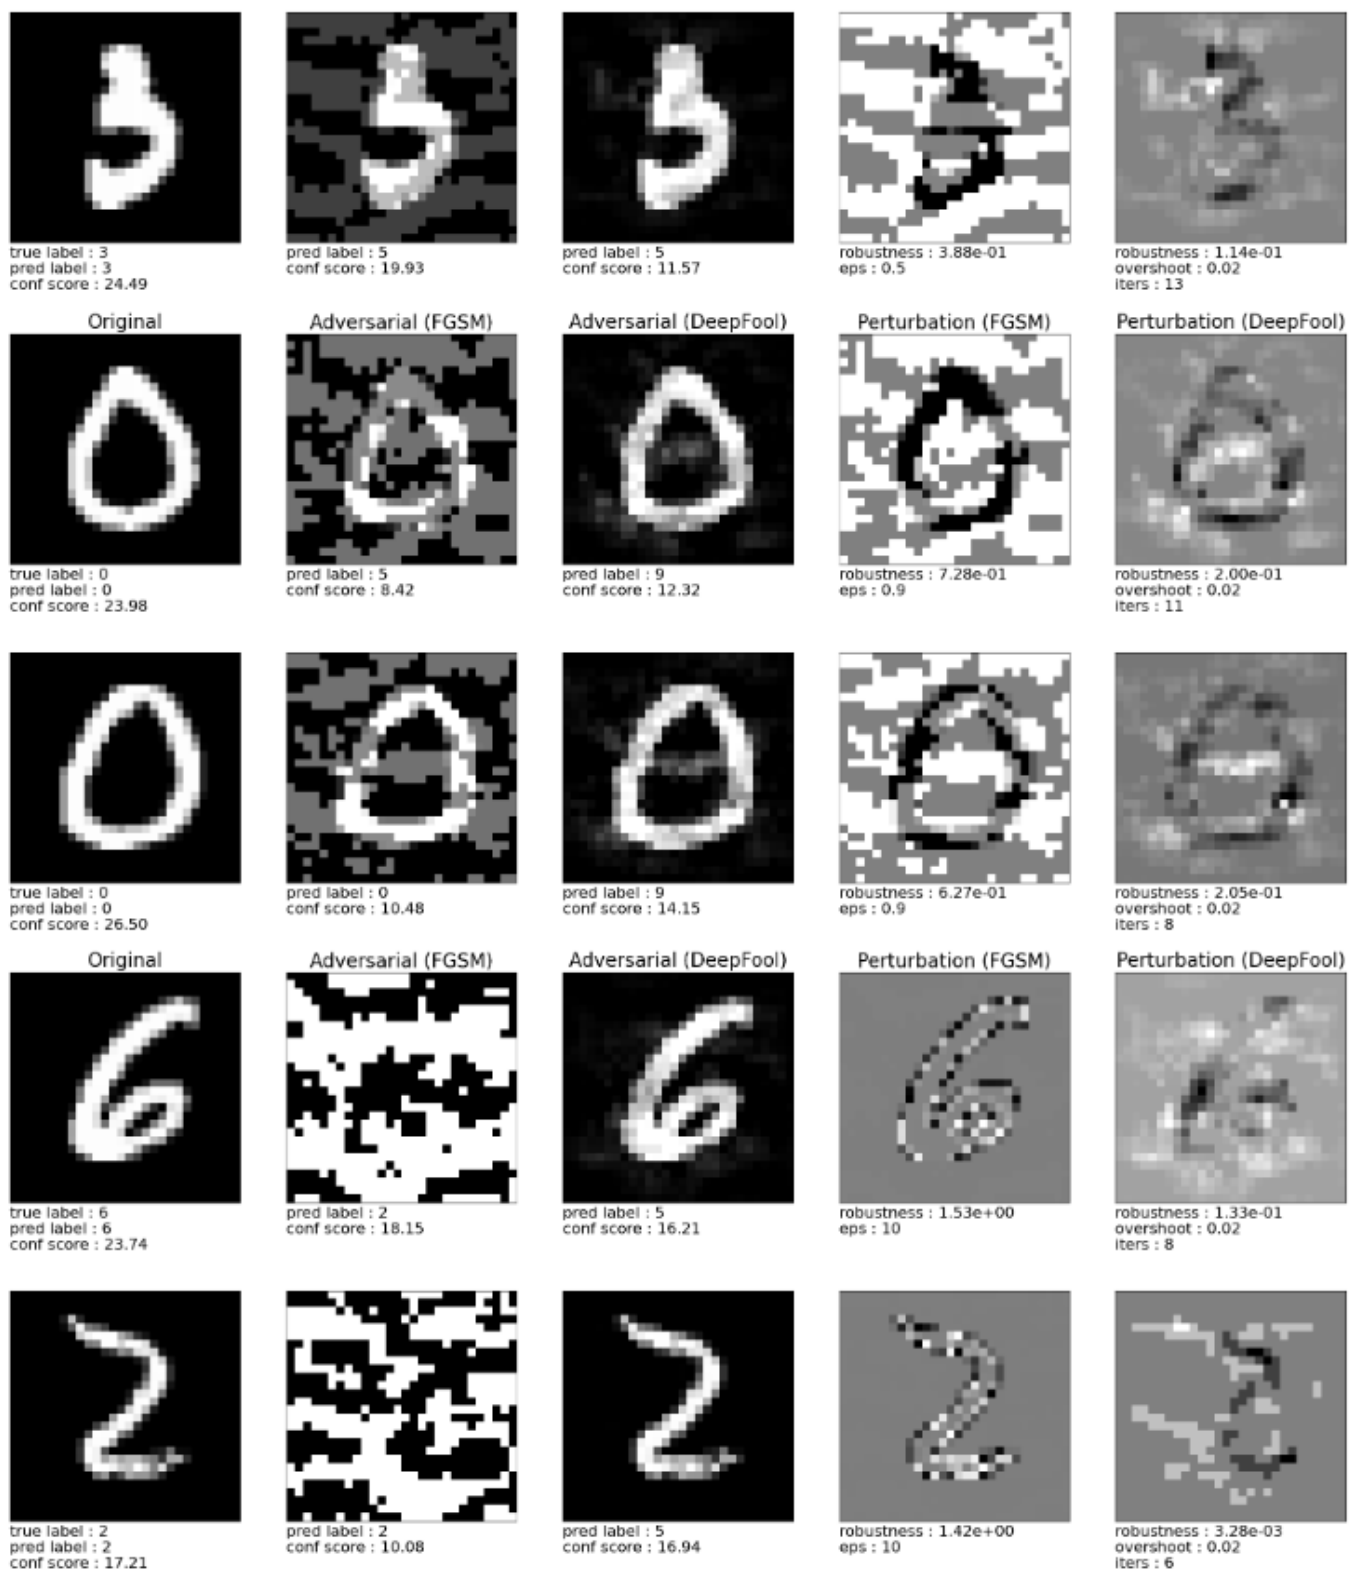
fgsm_eps = 0.02
model = LeNet_MNIST().to(device)
model.load_state_dict(torch.load('weights/clean/mnist_lenet.pth'))
display_attack(device, model, mnist_test, mnist_tf_inv, mnist_min, mnist_max, fgsm_eps, deep_args, has_labels=False, l2_norm=0)
if device.type == 'cuda': torch.cuda.empty_cache()

fgsm_eps = 0.5
model = LeNet_MNIST().to(device)
model.load_state_dict(torch.load('weights/clean/mnist_lenet.pth'))
display_attack(device, model, mnist_test, mnist_tf_inv, mnist_min, mnist_max, fgsm_eps, deep_args, has_labels=False, l2_norm=0)
if device.type == 'cuda': torch.cuda.empty_cache()

fgsm_eps = 0.9
model = LeNet_MNIST().to(device)
model.load_state_dict(torch.load('weights/clean/mnist_lenet.pth'))
display_attack(device, model, mnist_test, mnist_tf_inv, mnist_min, mnist_max, fgsm_eps, deep_args, has_labels=False, l2_norm=0)
if device.type == 'cuda': torch.cuda.empty_cache()

fgsm_eps = 10
model = LeNet_MNIST().to(device)
model.load_state_dict(torch.load('weights/clean/mnist_lenet.pth'))
display_attack(device, model, mnist_test, mnist_tf_inv, mnist_min, mnist_max, fgsm_eps, deep_args, has_labels=False, l2_norm=0)
if device.type == 'cuda': torch.cuda.empty_cache()
```

Original	Adversarial (FGSM)	Adversarial (DeepFool)	Perturbation (FGSM)	Perturbation (DeepFool)
 true label : 4 pred label : 4 conf score : 29.95	 pred label : 4 conf score : 29.95	 pred label : 9 conf score : 19.23	 robustness : 8.12e-04 eps : 0.001	 robustness : 1.24e-01 overshoot : 0.02 iters : 10
 true label : 5 pred label : 5 conf score : 12.22	 pred label : 5 conf score : 12.15	 pred label : 3 conf score : 8.95	 robustness : 8.28e-04 eps : 0.001	 robustness : 3.36e-02 overshoot : 0.02 iters : 9
 true label : 7 pred label : 7 conf score : 15.81	 pred label : 7 conf score : 15.02	 pred label : 9 conf score : 11.57	 robustness : 1.57e-02 eps : 0.02	 robustness : 6.26e-02 overshoot : 0.02 iters : 10
 true label : 0 pred label : 0 conf score : 20.85	 pred label : 0 conf score : 20.36	 pred label : 9 conf score : 11.75	 robustness : 1.67e-02 eps : 0.02	 robustness : 1.91e-01 overshoot : 0.02 iters : 8
 true label : 4 pred label : 4 conf score : 28.11	 pred label : 9 conf score : 26.76	 pred label : 9 conf score : 18.26	 robustness : 3.80e-01 eps : 0.5	 robustness : 8.73e-02 overshoot : 0.02 iters : 7



15. Проверим влияние параметра `fgsm_eps` для NiN на датасете Cifar-10.

```

fgsm_eps = 0.001
model = Net().to(device)
model.load_state_dict(torch.load('weights/clean/cifar_nin.pth'))
display_attack(device, model, cifar_test, cifar_tf_inv, cifar_min, cifar_max, fgsm_eps, deep_args, has_labels=False, l2_norm=0)
if device.type == 'cuda': torch.cuda.empty_cache()

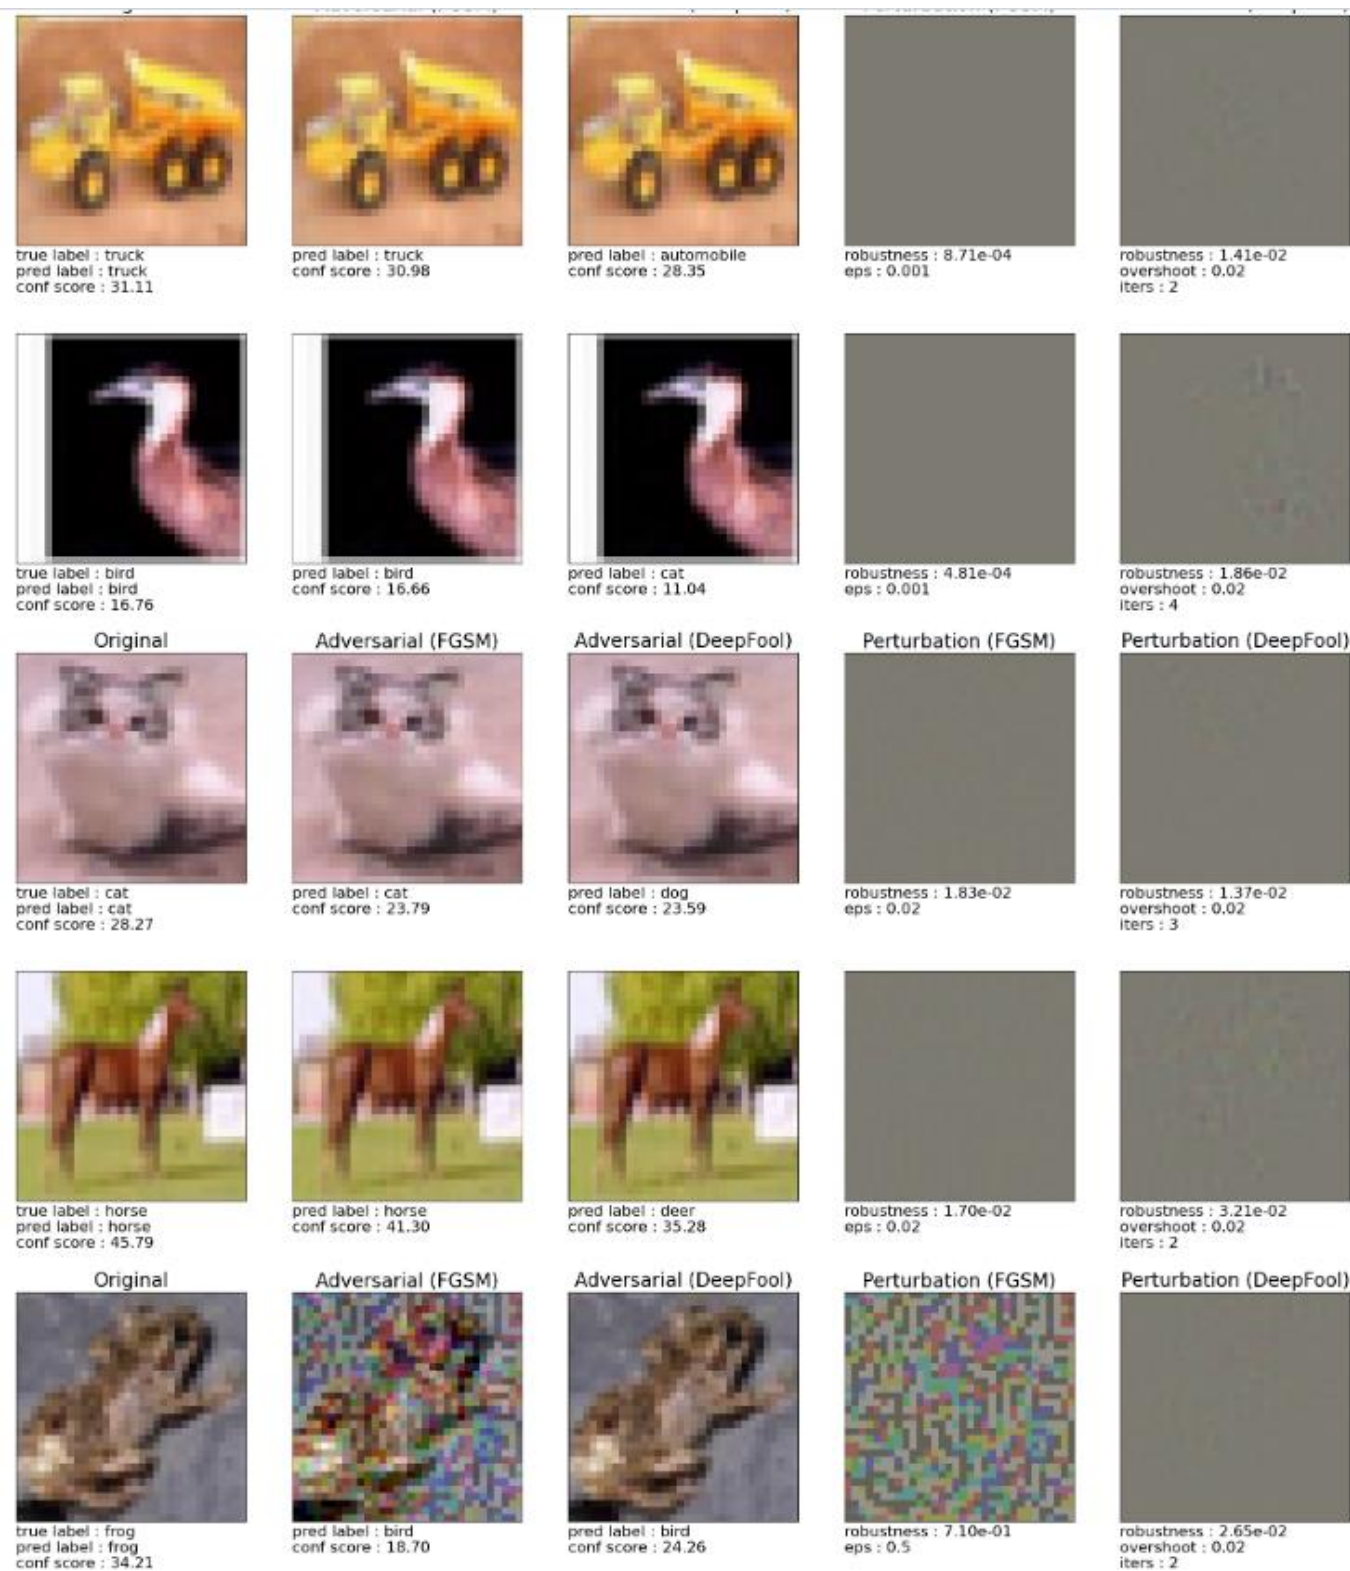
fgsm_eps = 0.02
model = Net().to(device)
model.load_state_dict(torch.load('weights/clean/cifar_nin.pth'))
display_attack(device, model, cifar_test, cifar_tf_inv, cifar_min, cifar_max, fgsm_eps, deep_args, has_labels=False, l2_norm=0)
if device.type == 'cuda': torch.cuda.empty_cache()

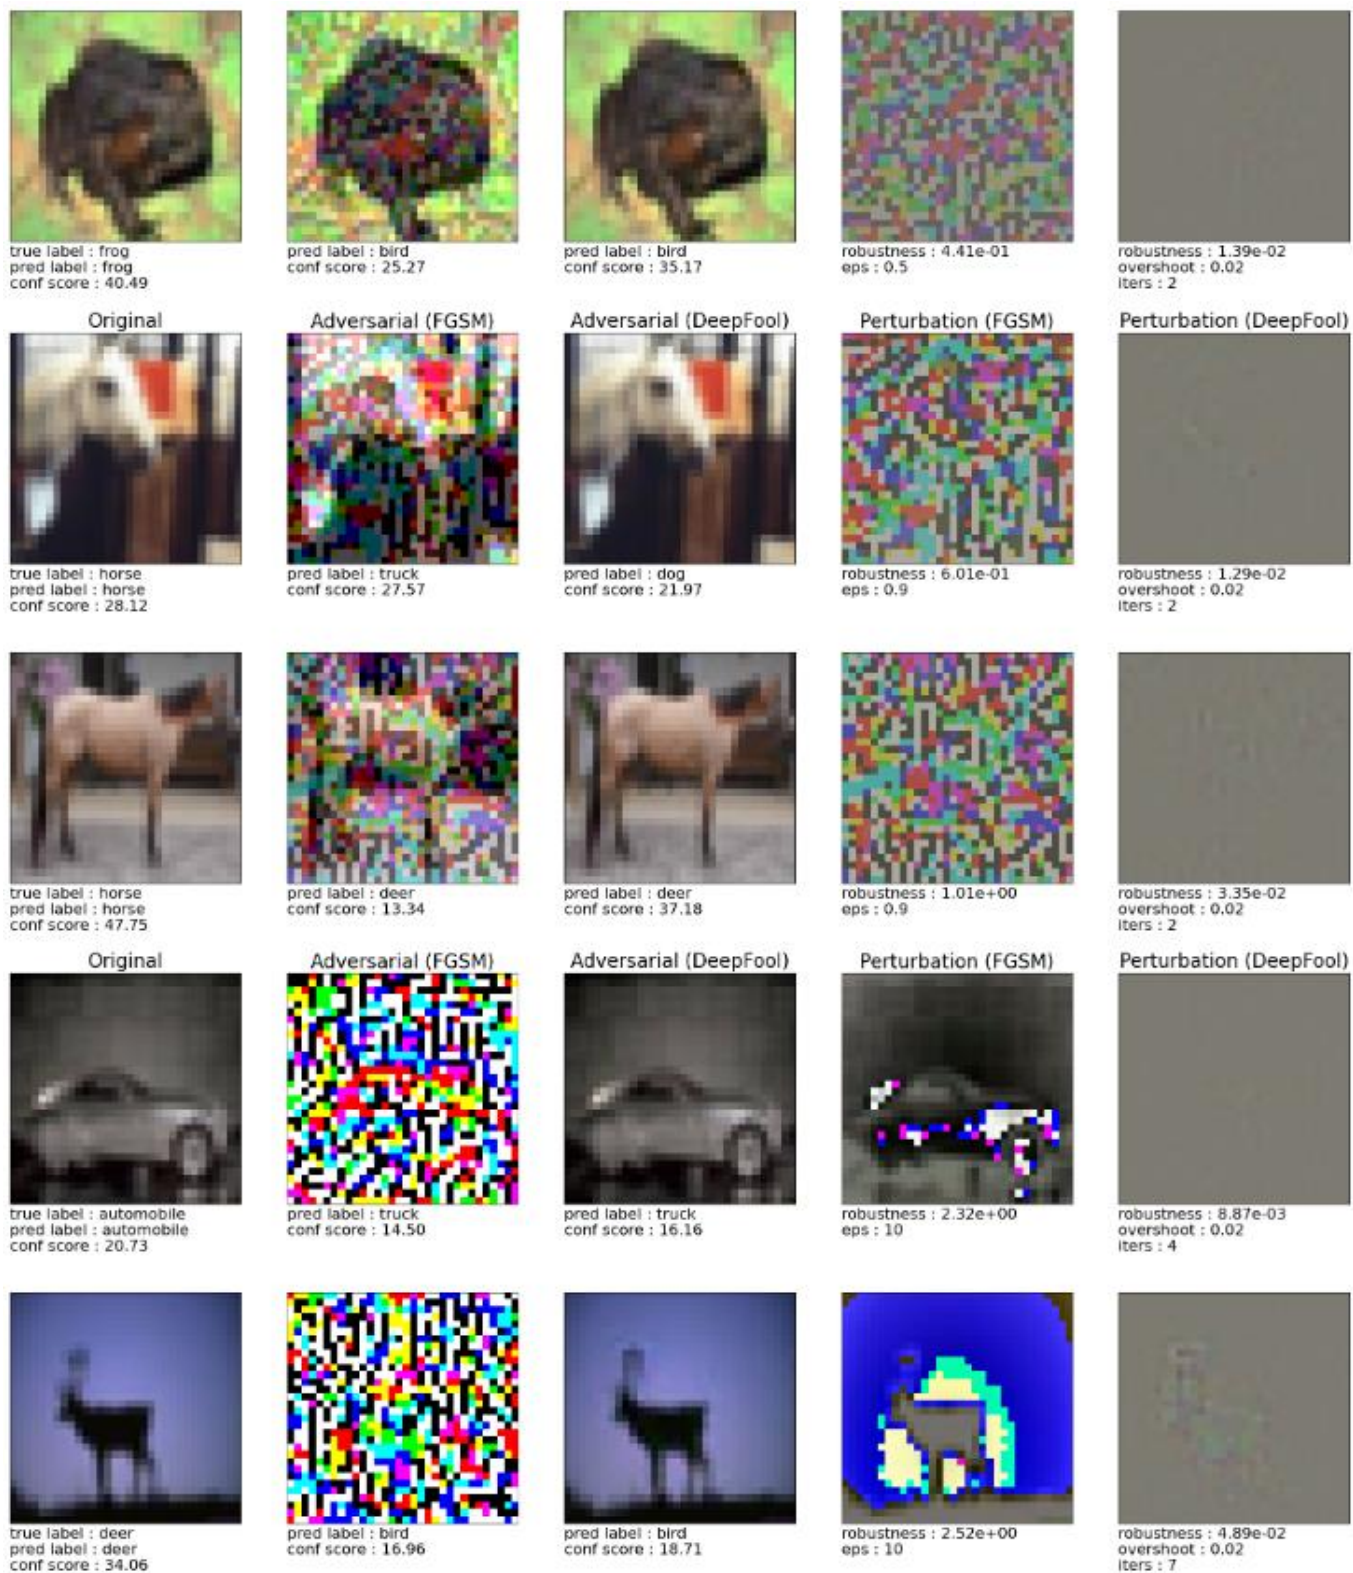
fgsm_eps = 0.5
model = Net().to(device)
model.load_state_dict(torch.load('weights/clean/cifar_nin.pth'))
display_attack(device, model, cifar_test, cifar_tf_inv, cifar_min, cifar_max, fgsm_eps, deep_args, has_labels=False, l2_norm=0)
if device.type == 'cuda': torch.cuda.empty_cache()

fgsm_eps = 0.9
model = Net().to(device)
model.load_state_dict(torch.load('weights/clean/cifar_nin.pth'))
display_attack(device, model, cifar_test, cifar_tf_inv, cifar_min, cifar_max, fgsm_eps, deep_args, has_labels=False, l2_norm=0)
if device.type == 'cuda': torch.cuda.empty_cache()

fgsm_eps = 10
model = Net().to(device)
model.load_state_dict(torch.load('weights/clean/cifar_nin.pth'))
display_attack(device, model, cifar_test, cifar_tf_inv, cifar_min, cifar_max, fgsm_eps, deep_args, has_labels=False, l2_norm=0)
if device.type == 'cuda': torch.cuda.empty_cache()

```





16. Проверим влияние параметра fgsm_eps для LeNet на датасете Cifar-10.

```

fgsm_eps = 0.001
model = LeNet_CIFAR().to(device)
model.load_state_dict(torch.load('weights/clean/cifar_lenet.pth'))
display_attack(device, model, cifar_test, cifar_tf_inv, cifar_min, cifar_max, fgsm_eps, deep_args, has_labels=False, l2_norm=0)
if device.type == 'cuda': torch.cuda.empty_cache()









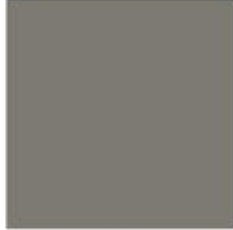












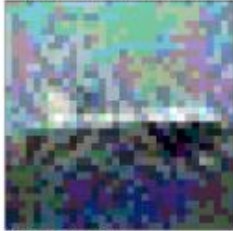

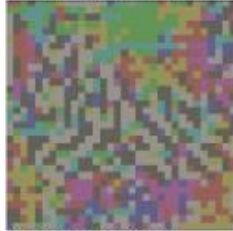

fgsm_eps = 0.02
model = LeNet_CIFAR().to(device)
model.load_state_dict(torch.load('weights/clean/cifar_lenet.pth'))
display_attack(device, model, cifar_test, cifar_tf_inv, cifar_min, cifar_max, fgsm_eps, deep_args, has_labels=False, l2_norm=0)
if device.type == 'cuda': torch.cuda.empty_cache()

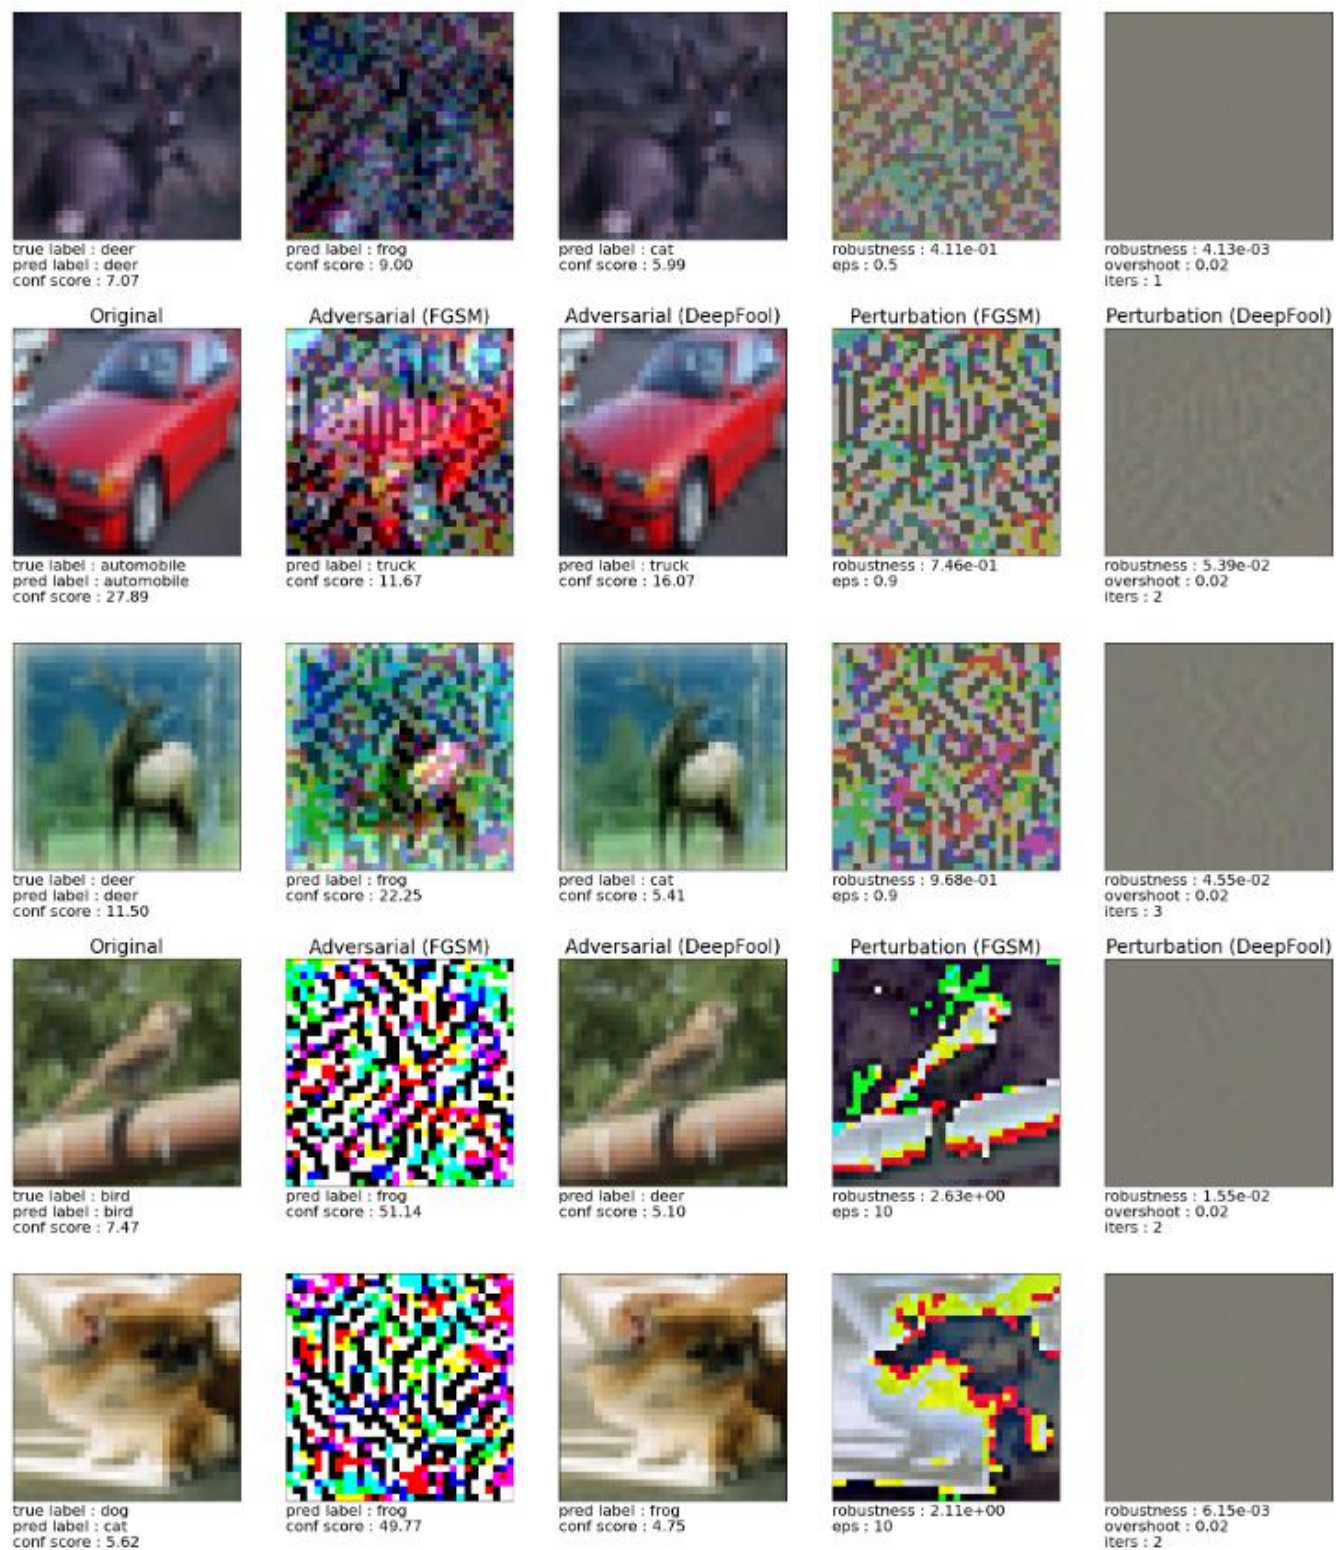
fgsm_eps = 0.5
model = LeNet_CIFAR().to(device)
model.load_state_dict(torch.load('weights/clean/cifar_lenet.pth'))
display_attack(device, model, cifar_test, cifar_tf_inv, cifar_min, cifar_max, fgsm_eps, deep_args, has_labels=False, l2_norm=0)
if device.type == 'cuda': torch.cuda.empty_cache()

fgsm_eps = 0.9
model = LeNet_CIFAR().to(device)
model.load_state_dict(torch.load('weights/clean/cifar_lenet.pth'))
display_attack(device, model, cifar_test, cifar_tf_inv, cifar_min, cifar_max, fgsm_eps, deep_args, has_labels=False, l2_norm=0)
if device.type == 'cuda': torch.cuda.empty_cache()

fgsm_eps = 10
model = LeNet_CIFAR().to(device)
model.load_state_dict(torch.load('weights/clean/cifar_lenet.pth'))
display_attack(device, model, cifar_test, cifar_tf_inv, cifar_min, cifar_max, fgsm_eps, deep_args, has_labels=False, l2_norm=0)
if device.type == 'cuda': torch.cuda.empty_cache()

```


Original	Adversarial (FGSM)	Adversarial (DeepFool)	Perturbation (FGSM)	Perturbation (DeepFool)
 true label : deer pred label : deer conf score : 8.35	 pred label : deer conf score : 8.29	 pred label : horse conf score : 7.16	 robustness : 8.19e-04 eps : 0.001	 robustness : 1.60e-02 overshoot : 0.02 iters : 3
 true label : truck pred label : automobile conf score : 7.19	 pred label : automobile conf score : 7.32	 pred label : ship conf score : 6.05	 robustness : 7.32e-04 eps : 0.001	 robustness : 4.48e-03 overshoot : 0.02 iters : 2
 true label : horse pred label : horse conf score : 24.45	 pred label : horse conf score : 21.08	 pred label : dog conf score : 13.60	 robustness : 1.83e-02 eps : 0.02	 robustness : 4.95e-02 overshoot : 0.02 iters : 3
 true label : truck pred label : cat conf score : 2.91	 pred label : cat conf score : 3.28	 pred label : dog conf score : 2.61	 robustness : 1.99e-02 eps : 0.02	 robustness : 5.30e-03 overshoot : 0.02 iters : 3
 true label : ship pred label : ship conf score : 17.76	 pred label : bird conf score : 4.53	 pred label : airplane conf score : 10.46	 robustness : 4.61e-01 eps : 0.5	 robustness : 3.11e-02 overshoot : 0.02 iters : 2



Вывод: параметр fgsm_eps влияет на устойчивость сети. При увеличении значения fgsm_eps сети становятся уязвимыми к атакам.